

COMS 4771  
Assignment Project Exam Help

Regression  
<https://powcoder.com>

Add WeChat powcoder

Nakul Verma

# Last time...

- Support Vector Machines
- Maximum Margin formulation
- Constrained Optimization
- Lagrange Duality Theory
- Convex Optimization
- SVM dual and Interpretation
- How get the optimal solution

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Learning more Sophisticated Outputs

So far we have focused on classification  $f: X \rightarrow \{1, \dots, k\}$

What about **other outputs**?


- $PM_{2.5}$  (pollutant) particulate matter exposure estimate:

**Input:** # cars, temperature, etc.

**Output:** 50 ppb

<https://powcoder.com>

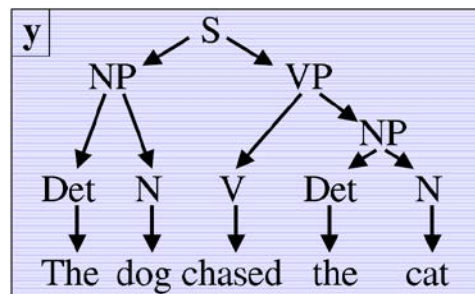
- Pose estimation



Add WeChat powcoder

- Sentence structure estimate:

**x** The dog chased the cat



# Regression

We'll focus on problems with real number outputs (regression problem):

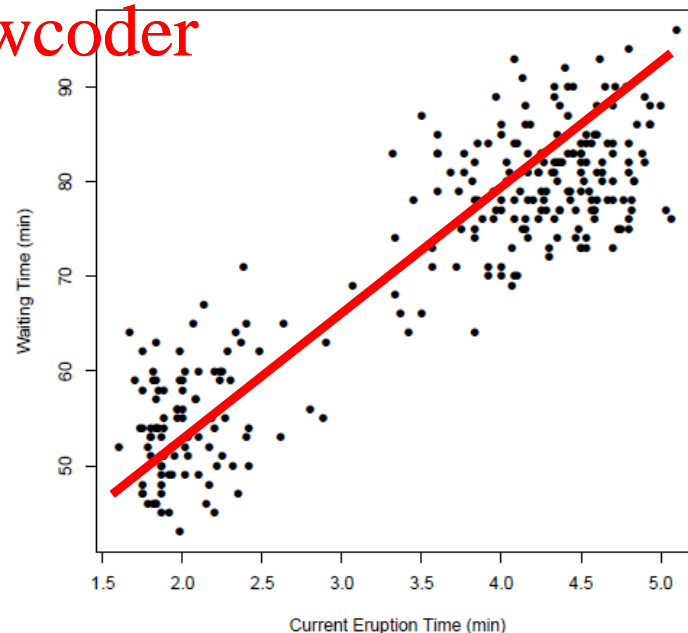
$$f : X \rightarrow \mathbb{R}$$

Example: **Assignment Project Exam Help**

**<https://powcoder.com>**

Next eruption time of old faithful geyser (at Yellowstone)

**Add WeChat powcoder**

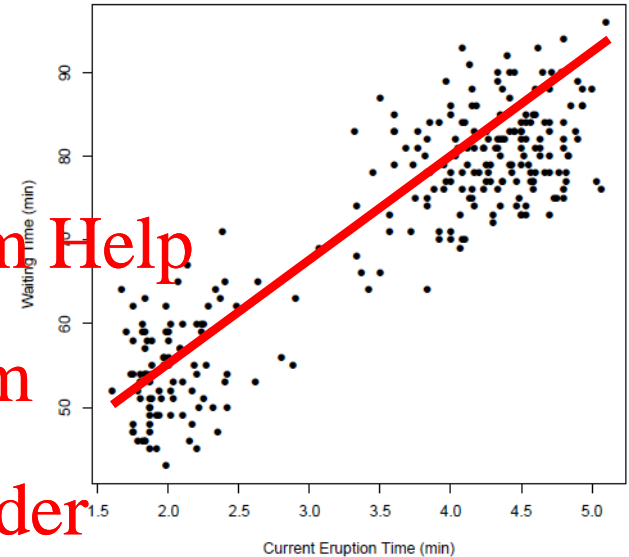


# Regression Formulation for the Example

Given  $x$ , want to predict an estimate  $\hat{y}$  of  $y$ , which minimizes the discrepancy ( $L$ ) between  $\hat{y}$  and  $y$ .

Loss  $\rightarrow$   $L(\hat{y}; y) := |\hat{y} - y|$  *Absolute error*  
 $L(\hat{y}; y) := (\hat{y} - y)^2$  *Squared error*

<https://powcoder.com>



A **linear predictor**  $f$ , can be defined by the slope  $w$  and the intercept  $w_0$ :

$$\hat{f}(\vec{x}) := \vec{w} \cdot \vec{x} + w_0$$

which minimizes the prediction loss.

$$\min_{w, w_0} \mathbb{E}_{\vec{x}, y} [L(\hat{f}(\vec{x}), y)]$$

*How is this different from **classification**?*


# Parametric vs non-parametric Regression

If we assume a particular form of the regressor:

*Parametric regression*

*Goal: to learn the parameters which yield the minimum error/loss*


<https://powcoder.com>



If no specific form of regressor is assumed.

*Non-parametric regression*

*Goal: to learn the predictor directly from the input data that yields the minimum error/loss*



# Linear Regression

Want to find a **linear predictor**  $f$ , i.e.,  $w$  (intercept  $w_0$  absorbed via lifting):

$$\hat{f}(\vec{x}) := \vec{w} \cdot \vec{x}$$

which minimizes the prediction loss over the population.


<https://powcoder.com>

Add WeChat powcoder

We estimate the parameters by minimizing the corresponding loss on the training data:

$$\begin{aligned} \arg \min_w \frac{1}{n} \sum_{i=1}^n [L(\vec{w} \cdot \vec{x}_i, y_i)] \\ = \arg \min_w \frac{1}{n} \sum_{i=1}^n (\vec{w} \cdot \vec{x}_i - y_i)^2 \end{aligned}$$

*for squared error*



# Linear Regression: Learning the Parameters

Linear predictor with squared loss:

$$\arg \min_w \frac{1}{n} \sum_{i=1}^n (\vec{w} \cdot \vec{x}_i - y_i)^2$$

Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

$$= \arg \min_w \left\| \begin{pmatrix} \dots x_1 \dots \\ \dots x_i \dots \\ \dots x_n \dots \end{pmatrix} \begin{pmatrix} w \end{pmatrix} - \begin{pmatrix} y_1 \\ y_i \\ y_n \end{pmatrix} \right\|_2^2$$

$$= \arg \min_w \|X\vec{w} - \vec{y}\|_2^2$$

*Unconstrained  
problem!*

*Can take the gradient and  
examine the stationary points!*

*Why need not check the  
second order conditions?*



# Linear Regression: Learning the Parameters

Best fitting  $w$ :

$$\frac{\partial}{\partial \vec{w}} \|X\vec{w} - \vec{y}\|^2 = 2X^T(X\vec{w} - \vec{y})$$

$$X^T X \vec{w} = X^T \vec{y}$$

$$\Rightarrow \vec{w}_{\text{ols}} = (X^T X)^{\dagger} X^T \vec{y}$$

Pseudo-inverse

*The solution is unique and stable when  $X^T X$  is invertible*

*What is the interpretation of this solution?*

# Linear Regression: Geometric Viewpoint

Consider the **column space** view of data  $\mathbf{X}$ :

$$\begin{pmatrix} \vdots & \mathbf{x}_1 & \vdots \\ \vdots & \mathbf{x}_i & \vdots \\ \vdots & \mathbf{x}_n & \vdots \end{pmatrix} \quad \ddot{x}_1, \dots, \ddot{x}_d \in \mathbf{R}^n$$

Assignment Project Exam Help

Find a  $\mathbf{w}$ , such that the linear combination of **minimizes**

$$\frac{1}{n} \left\| \vec{y} - \sum_{i=1}^d w_i \ddot{x}_i \right\|^2 =: \text{residual} \quad y = X \vec{w}_{\text{ols}} = X(X^\top X)^\dagger X^\top \vec{y}$$

https://powcoder.com  
Add WeChat powcoder


projection matrix  $\Pi$

Say  $\hat{y}$  is the ols solution, ie,

$$\hat{y} := X \vec{w}_{\text{ols}} = \sum_{i=1}^d w_{\text{ols},i} \ddot{x}_i$$

Thus,  $\hat{y}$  is the **orthogonal projection** of  $y$  onto the  $\text{span}(\ddot{x}_1, \dots, \ddot{x}_d)$  !

$w_{\text{ols}}$  forms the **coefficients** of  $\hat{y}$



# Linear Regression: Statistical Modeling View

Let's assume that data is **generated** from the following process:

- A example  $x_i$  is draw independently from the data space  $\mathbf{X}$

$x_i \sim \mathcal{D}_X$   
Assignment Project Exam Help

- $y_{\text{clean}}$  is computed as  $(w \cdot x_i)$ , from a fixed unknown  $w$

$y_{\text{clean}} := w \cdot x_i$   
Add WeChat powcoder

- $y_{\text{clean}}$  is corrupted from by adding independent Gaussian noise  $N(0, \sigma^2)$

$$y_i := y_{\text{clean}} + \epsilon_i = w \cdot x_i + \epsilon_i \quad \epsilon_i \sim N(0, \sigma^2)$$

- $(x_i, y_i)$  is revealed as the  $i^{\text{th}}$  sample

$$(x_1, y_1), \dots, (x_n, y_n) =: S$$

# Linear Regression: Statistical Modeling View

How can we determine  $w$ , from Gaussian noise corrupted observations?

$$S = (x_1, y_1), \dots, (x_n, y_n)$$

Observation:

**Assignment Project Exam Help**

$$y_i \sim w \cdot x_i + N(0, \sigma^2) = N(w \cdot x_i, \sigma^2)$$

**<https://powcoder.com>**

*How to estimate  
parameters of a Gaussian?*

parameter

$$\begin{aligned} \log \mathcal{L}(w|S) &= \sum_{i=1}^n \log p(y_i|w) \\ &\propto \sum_{i=1}^n \frac{-(w \cdot x_i - y_i)^2}{2\sigma^2} \end{aligned}$$

**Add WeChat powcoder**

*Let's try Maximum  
Likelihood Estimation!*

*ignoring terms  
independent of  $w$*



*optimizing for  $w$  yields  
the same ols result!*

*What happens if we model each  $y_i$  with  
indep. noise of different variance?*

# Linear Regression for Classification?

Linear regression seems general, can we use it to derive a binary classifier?

Let's study 1-d data:



*Problem #2: Not really linear!*

*Perhaps it is linear in some  
transformed coordinates?*

# Linear Regression for Classification



Interpretation:

For an event that occurs with probability  $P$ , the **odds** of that event is:

$$\text{odds}(P) := \frac{P}{1 - P}$$

*For an event with  $P=0.9$ , odds = 9*

*But, for an event  $P=0.1$ , odds = 0.11*

**(very asymmetric)**

Consider the “log” of the odds

$$\log(\text{odds}(P)) := \text{logit}(P) := \log\left(\frac{P}{1 - P}\right)$$

$$\text{logit}(P) = -\text{logit}(1 - P)$$

**Symmetric!**

# Logistic Regression

Model the log-odds or logit with linear function!


$$\text{logit}(P(x)) = \log \left( \frac{P(x)}{1 - P(x)} \right) = w \cdot x$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

OK, we have a model, how do we learn the parameters?



# Logistic Regression: Learning Parameters

Given samples  $S = (x_1, y_1), \dots, (x_n, y_n)$  ( $y_i \in \{0,1\}$  binary)

$$\mathcal{L}(w|S) = \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1-y_i} \quad \text{Binomial}$$

$$\log \mathcal{L}(w|S) = \sum_{i=1}^n y_i \log p(x_i) + (1 - y_i) \log(1 - p(x_i))$$

$$= \sum_{i=1}^n \log 1 - p(x_i) + \sum_{i=1}^n y_i \log \frac{p(x_i)}{1 - p(x_i)}$$

$$= \sum_{i=1}^n -\log 1 + e^{w \cdot x_i} + \sum_{i=1}^n y_i w \cdot x_i$$

*Can take the derivative and analyze stationary points,  
unfortunately no closed form solution*

*(use iterative methods like gradient descent to find the solution)*

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Now, use logistic model!



# Linear Regression: Other Variations

Back to the ordinary least squares (ols):

$$\text{minimize } \|X\vec{w} - \vec{y}\|_2^2$$

**Assignment Project Exam Help**

*Often poorly behaved  
when  $X^T X$  not invertible*

<https://powcoder.com>

Additionally how can we incorporate prior knowledge?

- perhaps want  $w$  to be sparse. **Add WeChat powcoder**  
*Lasso regression*
- perhaps want to simple  $w$ .  
*Ridge regression*

# Ridge Regression

## Objective

$$\text{minimize } \|X\vec{w} - \vec{y}\|^2 + \lambda \|\vec{w}\|^2$$

reconstruction error 'regularization' parameter

$$\vec{w}_{\text{ridge}} = (X^T X + \lambda I)^{-1} X^T \vec{y}$$

The 'regularization' helps avoid overfitting, and always resulting in a unique solution.


Equivalent to the following optimization problem:

$$\text{minimize } \|X\vec{w} - \vec{y}\|^2$$

$$\text{such that } \|\vec{w}\|^2 \leq B$$

Why?

Geometrically:



# Lasso Regression

## Objective

$$\text{minimize } \|X\vec{w} - \vec{y}\|^2 + \lambda \|\vec{w}\|_1$$

Assignment Project Exam Help

$\vec{w}_{\text{lasso}} = ?$  no closed form solution  
<https://powcoder.com>

Lasso regularization encourages sparse solutions.

Add WeChat powcoder

Equivalent to the following optimization problem:


$$\text{minimize } \|X\vec{w} - \vec{y}\|^2$$

$$\text{such that } \|\vec{w}\|_1 \leq B$$

Why?

How can we find the solution?

Geometrically:



# What About Optimality?

Linear regression (and variants) is great, but what can we say about the best possible estimate?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

*Can we construct an estimator for real outputs that  
**parallels** Bayes classifier for discrete outputs?*

# Optimal $L_2$ Regressor

Best possible regression estimate at  $x$ :  $f^*(x) := \mathbb{E}[Y|X = x]$

**Theorem:** for any regression estimate  $g(x)$

$$\mathbb{E}_{(x,y)} |f^*(x) - y|^2 \leq \mathbb{E}_{(x,y)} |g(x) - y|^2$$

*Assignment Project Exam Help* *Similar to Bayes classifier, but for regression.*

<https://powcoder.com>

Add WeChat powcoder

*Proof is straightforward...*

# Proof

Consider  $L_2$  error of  $g(x)$

$$f^*(x) := \mathbb{E}[Y|X = x]$$

$$\begin{aligned}\mathbb{E}|g(x) - y|^2 &= \mathbb{E}|g(x) - f^*(x) + f^*(x) - y|^2 \\ &= \mathbb{E}|g(x) - f^*(x)|^2 + \mathbb{E}|f^*(x) - y|^2\end{aligned}$$

*Why?*

Assignment Project Exam Help

**Cross term:**

$$\begin{aligned}2\mathbb{E}[(g(x) - f^*(x))(f^*(x) - y)] \\ &= 2\mathbb{E}_x[\mathbb{E}_{y|x}[(g(x) - f^*(x))(f^*(x) - y) | X = x]] \\ &= 2\mathbb{E}_x[(g(x) - f^*(x))\mathbb{E}_{y|x}[f^*(x) - y | X = x]] \\ &= 2\mathbb{E}_x[(g(x) - f^*(x))(f^*(x) - f^*(x))] = 0\end{aligned}$$

Therefore 
$$\mathbb{E}|g(x) - y|^2 = \int_x |g(x) - f^*(x)|^2 \mu(dx) + \mathbb{E}|f^*(x) - y|^2$$

*Which is minimized when  $g(x) = f^*(x)$ !*




# Non-parametric Regression

Linear regression (and variants) is great, but what if we don't know parametric form of the relationship between the independent and dependent variables?

How can we predict value of a new test point  $x$  **without** model assumptions?

Idea:

$\hat{y} = f(x) =$  *Average estimate  $\hat{Y}$  of observed data in a local neighborhood  $X$  of  $x$ !*



# Kernel Regression

$$\hat{y} = \hat{f}_n(x) := \sum_{i=1}^n w_i(x) y_i$$

Want weights that emphasize  
local observations

Assignment Project Exam Help

<https://powcoder.com>

Consider example localization functions:

$$K_h(x, x') = e^{-\|x - x'\|^2 / h}$$

Gaussian kernel

$$= \mathbf{1}[\|x - x'\| \leq h]$$

Box kernel

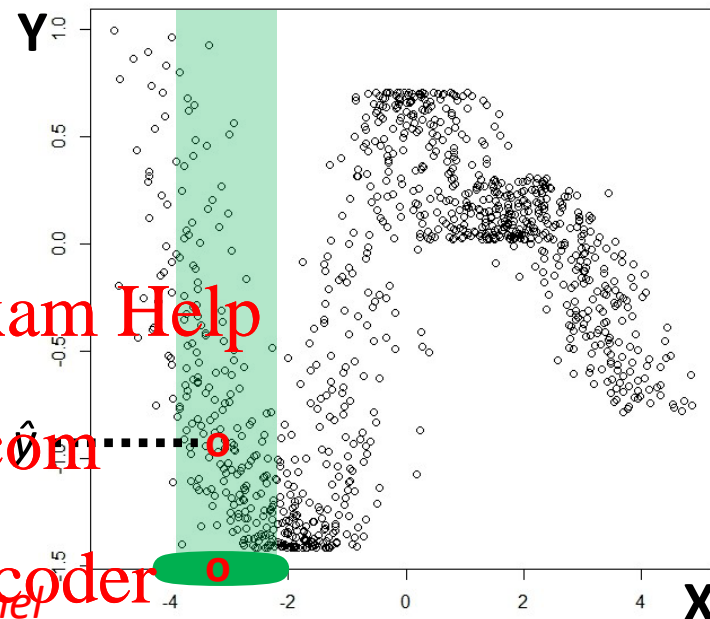
$$= [1 - (1/h)\|x - x'\|]_+$$

Triangle kernel

Then define:

$$w_i(x) := \frac{K_h(x, x_i)}{\sum_{j=1}^n K_h(x, x_j)}$$

Weighted average





# Consistency Theorem

Recall: best possible regression estimate at  $x$ :  $f^*(x) := \mathbb{E}[Y|X = x]$

## Assignment Project Exam Help

**Theorem:** As  $n \rightarrow \infty$ ,  $h \rightarrow 0$ ,  $hn \rightarrow \infty$ , then

<https://powcoder.com>

$\mathbb{E}_{(x,y)} |\hat{f}_{n,h}(x) - f^*(x)|^2 \rightarrow 0$   
Add WeChat powcoder

where  $\hat{f}_{n,h}(x) := \sum_{i=1}^n \frac{K_h(x, x_i)}{\sum_{j=1}^n K_h(x, x_j)} y_i$  is the kernel regressor with

most localization kernels.

*Proof is a bit tedious...*

# Proof Sketch

Prove for a fixed  $x$  and then integrate over (just like before)

$$\mathbb{E}|\hat{f}_{n,h}(x) - f^*(x)|^2 = \underbrace{\mathbb{E}[\hat{f}_{n,h}(x) - f^*(x)]^2}_{\text{squared bias of } \hat{f}_{n,h}} + \underbrace{\mathbb{E}[\hat{f}_{n,h}(x) - \mathbb{E}\hat{f}_{n,h}(x)]^2}_{\text{variance of } \hat{f}_{n,h}}$$

Assignment Project Exam Help

*Bias-variance  
decomposition*

<https://powcoder.com>

Sq. bias  $\approx c_1 h^2$

Add WeChat powcoder

Variance  $\approx c_2 \frac{1}{nh^d}$

Pick

$h \approx n^{-1/2+d}$

$$\mathbb{E}|\hat{f}_{n,h}(x) - f^*(x)|^2 \approx n^{-2/2+d} \rightarrow 0$$



# Kernel Regression

$$\hat{y} = \hat{f}_n(x) := \sum_{i=1}^n \frac{K_h(x, x_i)}{\sum_{j=1}^n K_h(x, x_j)} y_i$$

## Assignment Project Exam Help

### Advantages:

- Does not assume any parametric form of the regression function.
- Kernel regression is consistent.

<https://powcoder.com>

Add WeChat powcoder

### Disadvantages:

- Evaluation time complexity:  $O(dn)$
- Need to keep all the data around!


*How can we address the shortcomings of kernel regression?*

# kd trees: Speed Up Nonparametric Regression

k-d trees to the rescue!

Idea: partition the data in cells organized in a tree based hierarchy. (just like before)

To return an estimated value, return the average  $y$  value in a cell!



# What We Learned...

- Linear Regression
- Parametric vs Nonparametric regression  
**Assignment Project Exam Help**
- Logistic Regression for classification  
**<https://powcoder.com>**
- Ridge and Lasso Regression  
**Add WeChat powcoder**
- Kernel Regression
- Consistency of Kernel Regression
- Speeding non-parametric regression with trees

# Questions?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Next time...

Statistical Theory of Learning! Assignment! Project Exam Help

<https://powcoder.com>

Add WeChat powcoder